

A Bayesian Hierarchical Model For Predicting Song Popularity

Statistics 775

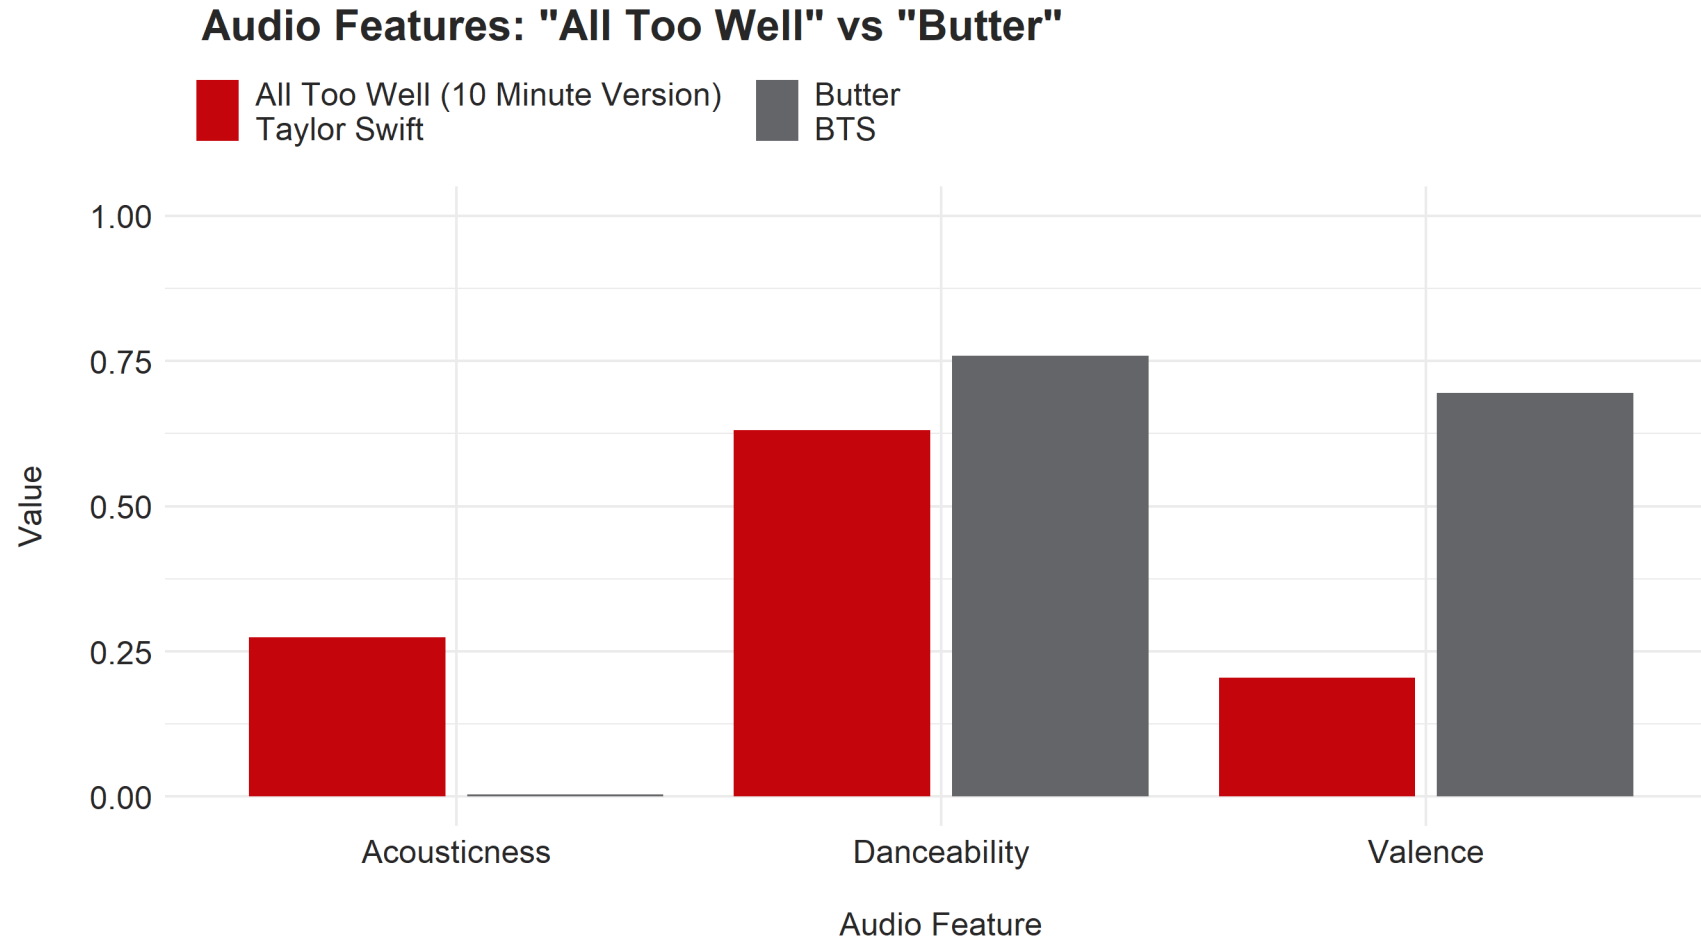
James Hubbs



Predicting Song Popularity From Audio Features

- Do things like tempo, rhythm, timbre, and other qualities of the sound help drive popularity?
- The main idea:
 - Use Spotify's Web API to sample 100 songs per year from 1970 – 2021 (so $n = 5100$)
 - For each song, extract its popularity and audio features (loudness, danceability, speechiness, valence, acousticness, and more)
 - Build a hierarchical model

Audio Features Are (Mostly) Intuitive




Audio Features Have Prominent Trends




Average Speechiness By Year

1970-2021



Average Acousticness By Year

1970-2021



Modeling Popularity

- Consider 10 predictors (audio features) and 6 groups (decades)
- Use a standard multiple linear regression setup (layer 1)
 - $Y_{it} \mid \alpha_t \beta_t \sigma^2 \sim N(\mu_{it}, \sigma^2)$ where $\mu_{it} = \alpha_t + x_i^T \beta_t$
- Allow intercept and slopes to vary by decade (layer 2)
 - $\alpha_t \mid \mu_\alpha \tau_\alpha^2 \sim N(\mu_\alpha, \tau_\alpha^2)$
 - $\beta_{jt} \mid \mu_{\beta_j} \tau_\beta^2 \sim N(\mu_{\beta_j}, \tau_\beta^2)$
- Set priors/hyperparameters using both prior predictive checks and prior knowledge regarding coefficient directions (layer 3)

Model Performance and Conclusions

- In general, audio features are only weakly predictive of popularity
 - The model suggests danceability is most influential
- Our perception of music is complicated!
- Spotify's audio features (the publicly-available ones, at least) don't capture everything